

High-Level Design (HLD)

AI-Powered Security Analytics Platform

Field	Value
Document ID	HLD-001
Version	1.0
Date	2026-07-18
Status	Draft
Architecture Reference	[ADR-001](file:///d:/AI%20SIEM/docs/architecture/decisions/ADR-001-modular-monolith.md)
Requirements Reference	[SRS-001](file:///d:/AI%20SIEM/docs/SRS.md)

Table of Contents

- 1. [Overall Architecture](#)
- 2. [Component Diagram](#)
- 3. [Data Flow](#)
- 4. [Module Responsibilities](#)
- 5. [Technology Stack](#)
- 6. [Communication Between Modules](#)

1. Overall Architecture

The AI-Powered Security Analytics Platform follows a **Modular Monolith** architecture. The system is composed of **4 independent processes** that communicate through **file-based I/O, HTTP REST**, and


WebSocket. There are no distributed message brokers, no service meshes, and no container orchestration.

1.1 Architectural Style

Property	Value
Style	Modular Monolith
Design Principles	Clean Architecture, SOLID, Repository Pattern, Dependency Injection
Processes	4 (Collector, Backend, AI Engine, Frontend)
Databases	3 (PostgreSQL, MongoDB, Redis)
Deployment	Docker Compose (single command)
Tenancy	Single-tenant (MVP)






1.2 System Context Diagram

This diagram shows the platform in context with its external actors and the environment it operates within.



1.3 High-Level Architecture Diagram


This diagram details the internal structure of each process and how they interconnect.



2. Component Diagram


2.1 Process-Level Component Diagram

This diagram shows the four processes and the three data stores as top-level components, with their interfaces and dependencies.




2.2 Backend Internal Component Diagram


The backend monolith is organized into **9 modules**, a **workers layer**, and a **shared kernel**. Each module follows Clean Architecture with 4 layers: Domain, Application, Infrastructure, Interface.




2.3 AI Engine Component Diagram




2.4 Collector Component Diagram



2.5 Frontend Component Diagram







3. Data Flow


3.1 End-to-End Data Flow

This diagram traces a security event from its source through the entire platform to the analyst's screen.







3.2 Collector Internal Data Flow




3.3 Backend Pipeline Data Flow (Worker Detail)






3.4 Authentication & Authorization Flow



3.5 Real-Time Dashboard Update Flow



4. Module Responsibilities

4.1 Collector (C++) — Process 1

Module	Responsibility	Inputs	Outputs
Syslog Receiver	Listen on UDP/TCP port 514, receive RFC 5424 syslog messages	Network packets	Raw syslog strings
File Tailer	Monitor log files for new lines, track read position	File system	Raw log lines
ETW Reader	Subscribe to Windows Event Log channels via ETW	Windows Event Tracing	Structured event records
HTTP Receiver	Accept JSON-formatted log events via REST endpoint	HTTP POST requests	JSON objects
OCSF Converter	Map source-specific fields to OCSF schema; place unknown fields in unmapped	Raw events from any source	OCSF JSON objects

Module	Responsibility	Inputs	Outputs
Batcher	Accumulate events until batch threshold (count or time) is reached	Individual OCSF events	Array of OCSF events
Atomic Writer	Write batch to <code>.tmp</code> file, then rename to <code>.json</code> for atomic visibility	Batch array	File in collector directory
Checkpoint Manager	Persist file read positions and sequence numbers for crash recovery	Read positions from sources	Checkpoint file on disk
Heartbeat Writer	Periodically write health status (timestamp, CPU, memory, events collected)	Internal metrics	Heartbeat file in collector directory

4.2 Backend (Node.js Express) — Process 2

4.2.1 Ingestion Module

Component	Responsibility
Directory Watcher	Monitors collector directory using <code>chokidar</code> ; detects new <code>.json</code> files; enqueues jobs into <code>IProcessingQueue</code>
IProcessingQueue	Abstract interface for the processing queue. Methods: <code>enqueue()</code> , <code>dequeue()</code> , <code>retry()</code> , <code>deadLetter()</code> , <code>getStatus()</code>
BullMQQueue	Infrastructure implementation of <code>IProcessingQueue</code> using BullMQ (Redis-backed). Handles persistence, retry logic, priority, and dead-letter

4.2.2 Workers

Component	Responsibility
Pipeline Worker	Dequeues jobs from <code>IProcessingQueue</code> independently of Directory Watcher. Orchestrates the full processing pipeline: <code>parse</code> → <code>normalize</code> → <code>extract features</code> → <code>rule engine + AI</code> → <code>correlate</code> → <code>score</code> → <code>store</code> . Multiple workers run concurrently (configurable: 1–16)

4.2.3 Parsing Module

Component	Responsibility
Parser	Reads batch file content, parses JSON, validates structure. Detects format (OCSF JSON expected from collector). Routes malformed files to quarantine
Normalizer	Validates each parsed event against the OCSF JSON Schema. Flags schema violations. Applies enrichment mappings (e.g., IP → GeoIP if available)

4.2.4 Analysis Module

Component	Responsibility
Feature Extractor	Computes statistical and contextual features from normalized events: IP entropy, login frequency, time-of-day deviation, byte volume anomaly, process execution rarity
Rule Engine	Evaluates feature-enriched events against Sigma-compatible detection rules loaded from PostgreSQL. Returns rule match alerts with severity and rule metadata
AI Client	HTTP client that sends feature vectors to the Python FastAPI AI Engine. Handles timeouts gracefully (fallback to rule-only). Implements <code>IAIClient</code> interface

4.2.5 Correlation Module

Component	Responsibility
Incident Correlator	Groups related alerts (from Rule Engine and AI Engine) by entity (user, host, IP), time window (configurable, default 15 min), and kill chain stage. Creates new incidents or merges alerts into existing open incidents
Risk Scorer	Computes composite risk score: <code>rule_weight × ML_confidence × asset_criticality</code> . Assigns severity level (Critical, High, Medium, Low, Informational). Handles missing factors with neutral defaults (1.0)

4.2.6 Incidents Module

Component	Responsibility
Incident Service	CRUD operations for incidents. Lifecycle management: Open → Investigating → Resolved → Closed. Assignment, notes, linking to constituent alerts and events

4.2.7 Rules Module

Component	Responsibility
Rule Service	CRUD operations for Sigma-compatible detection rules. Enable/disable without deletion. Import/export in YAML format. Rule testing against historical events (dry-run)

4.2.8 Auth Module

Component	Responsibility
Auth Service	Local username/password authentication with bcrypt. JWT token generation and validation (stored in Redis). RBAC enforcement (Admin, Security Engineer, SOC Analyst). Session management
Audit Logger	Records all user actions (login, rule changes, incident updates, config changes) as immutable audit entries in PostgreSQL

4.2.9 Collector Monitoring Module

Component	Responsibility
Collector Status Service	Reads heartbeat files from collector directory. Tracks: status (Online/Offline/Degraded), last heartbeat timestamp, files processed count, queue size (from <code>IProcessingQueue.getStatus()</code>), failed/quarantined file count, collector resource usage, last batch processed timestamp

4.2.10 Dashboard Module

Component	Responsibility
Dashboard Service	Aggregation queries for dashboard widgets: incident counts by severity, trend over time, top affected assets, recent incidents feed. Provides data for all dashboard visualizations

4.2.11 Shared Kernel

Component	Responsibility
OCSF Schema Definitions	TypeScript types/interfaces for all OCSF event categories used in the platform
Error Handling	Centralized error classes (DomainError, ApplicationError, InfrastructureError) with consistent structure
Logger	Structured JSON logging with correlation IDs, request tracing, and configurable log levels
Configuration	Environment-based configuration loading (12-Factor App). Validation on startup

4.3 AI Engine (Python FastAPI) — Process 3

Component	Responsibility
Anomaly Detection Endpoint	Receives feature vectors, runs Isolation Forest inference, returns anomaly scores
Explainability Endpoint	Computes SHAP values for a given prediction, returns feature importance rankings
Threat Classification Endpoint (optional)	Classifies events by threat category using Random Forest / XGBoost
Feature Preprocessor	Scales and encodes incoming features to match model training format
Model Loader	Loads serialized models (joblib/ONNX) at startup based on models/config.json version pointer

Component	Responsibility
Pydantic Schemas	Request/response validation for all endpoints

4.4 Frontend (Next.js) — Process 4

Page	Responsibility
Dashboard (/)	Real-time overview with incident counters, severity charts, timeline, top affected assets
Incidents (/incidents)	Incident list with filtering/sorting. Detail view with alerts, events, AI explanation, risk breakdown
Investigation (/investigate)	Log search with time range, source, severity, keyword filters. Event timeline. Pivot to related incidents
Rules (/rules)	Sigma rule YAML editor with syntax highlighting. Rule testing (dry-run). Import/export. Enable/disable
AI Insights (/ai-insights)	SHAP explanation visualizations. Model confidence scores. Anomaly score distributions
Collector Monitor (/collector)	Collector status, heartbeat, files processed, queue size, failed files, resource usage
Settings (/settings)	User management (CRUD + roles). Data source configuration. Retention policies
Audit Log (/audit)	Immutable record of all analyst and admin actions, filterable by user, action type, time range

5. Technology Stack

5.1 Complete Technology Map

Layer	Technology	Version Target	Purpose
Collector	C++	C++20	High-performance log collection and OCSF conversion
	CMake	3.20+	Cross-platform build system
	nlohmann/json	3.x	JSON serialization/deserialization
	yaml-cpp	0.7+	YAML configuration parsing
	spdlog	1.x	Fast structured logging
	Google Test	1.x	Unit testing framework
Backend	Node.js	20 LTS	JavaScript runtime
	Express	4.x	HTTP framework
	TypeScript	5.x	Type-safe JavaScript
	BullMQ	5.x	Redis-backed job queue (behind IProcessingQueue)
	chokidar	3.x	File system watcher
	pg (node-postgres)	8.x	PostgreSQL client
	Mongoose	8.x	MongoDB ODM
	ioredis	5.x	Redis client
	jsonwebtoken	9.x	JWT token management
	bcrypt	5.x	Password hashing
	Joi or Zod	latest	Request validation
	ws	8.x	WebSocket server
	Winston or Pino	latest	Structured logging

Layer	Technology	Version Target	Purpose
	Jest	29.x	Unit/integration testing
AI Engine	Python	3.11+	ML runtime
	FastAPI	0.100+	Async REST framework
	scikit-learn	1.x	Isolation Forest, Random Forest
	SHAP	0.x	Model explainability
	XGBoost	2.x	Optional threat classification
	Pydantic	2.x	Request/response validation
	joblib	1.x	Model serialization
	uvicorn	0.x	ASGI server
	pytest	8.x	Testing framework
Frontend	Next.js	14+	React SSR framework
	React	18+	UI library
	TypeScript	5.x	Type safety
	Recharts or Chart.js	latest	Dashboard charts
	Zustand or React Context	latest	State management
	Axios or Fetch	latest	REST API client
	Monaco Editor	latest	YAML rule editor
PostgreSQL	PostgreSQL	16	Relational database
MongoDB	MongoDB	7.x	Document database
Redis	Redis	7.x	Cache, queue, pub/sub
Deployment	Docker	latest	Containerization
	Docker Compose	v2	Multi-container orchestration

5.2 Technology Rationale Matrix

Decision	Chosen	Rejected Alternatives	Rationale
Collector language	C++	Go, Rust	Maximum throughput for log parsing; zero-copy potential; existing team skill
Backend framework	Express (Node.js)	Fastify, NestJS, Koa	Mature ecosystem, largest community, minimal magic, easy to reason about
AI framework	FastAPI (Python)	Flask, Django REST	Async-native, auto OpenAPI docs, Pydantic validation, best-in-class for ML serving
Frontend framework	Next.js	Vite + React, Remix	SSR for fast initial load, App Router, API routes for BFF, industry standard
Relational DB	PostgreSQL	MySQL, SQLite	JSONB support, row-level security, advanced indexing, proven at scale
Document DB	MongoDB	CouchDB, DynamoDB	Flexible schema for heterogeneous logs, TTL indexes, text search, strong Node.js ecosystem
Cache/Queue	Redis + BullMQ	In-memory queue, RabbitMQ	BullMQ provides persistent queue with retry/DLQ; Redis provides pub/sub + cache; single dependency


6. Communication Between Modules

6.1 Inter-Process Communication Matrix

From	To	Protocol	Data Format	Pattern	Latency
Collector	Backend	File system (atomic write + watch)	OCSF JSON files (.json)	Async, fire-and-forget	Seconds (batch interval)


From	To	Protocol	Data Format	Pattern	Latency
Collector	Backend	File system (heartbeat files)	JSON status file	Async, periodic	30s intervals
Backend	AI Engine	HTTP REST	JSON request/response	Sync, request- reply	< 200ms target
Frontend	Backend	HTTP REST	JSON request/response	Sync, request- reply	< 500ms target
Backend	Frontend	WebSocket (via Redis Pub/Sub)	JSON event push	Async, publish- subscribe	< 3s target

6.2 Inter-Process Communication Diagram



6.3 Intra-Backend Module Communication

Inside the backend monolith, modules communicate through **dependency injection** and **interfaces** — never through direct class instantiation or shared mutable state.



Note

Every diamond node (blue) in the diagram above is an **interface** defined in the domain layer. Concrete implementations live in the infrastructure layer and are injected at application startup. This is the **Dependency Inversion Principle** in action — the entire pipeline can be tested with mock implementations, and any component can be replaced without modifying its consumers.

6.4 API Contract Summary

Backend REST API (Express)

Method	Endpoint	Module	Description
POST	/api/v1/auth/login	Auth	Authenticate user, return JWT
POST	/api/v1/auth/logout	Auth	Invalidate session
GET	/api/v1/incidents	Incidents	List incidents (paginated, filterable)
GET	/api/v1/incidents/:id	Incidents	Get incident detail with alerts and events
PUT	/api/v1/incidents/:id	Incidents	Update incident (status, assignee, notes)
GET	/api/v1/events	Dashboard	Search normalized events (paginated, filterable)
GET	/api/v1/events/:id	Dashboard	Get single OCSF event detail
GET	/api/v1/rules	Rules	List detection rules
POST	/api/v1/rules	Rules	Create new Sigma rule
PUT	/api/v1/rules/:id	Rules	Update rule
DELETE	/api/v1/rules/:id	Rules	Delete rule

Method	Endpoint	Module	Description
POST	/api/v1/rules/:id/test	Rules	Dry-run rule against historical events
GET	/api/v1/dashboard/summary	Dashboard	Aggregated dashboard data
GET	/api/v1/collector/status	Collector Monitoring	Current collector health and metrics
GET	/api/v1/users	Auth	List users (Admin only)
POST	/api/v1/users	Auth	Create user (Admin only)
PUT	/api/v1/users/:id	Auth	Update user (Admin only)
GET	/api/v1/audit	Auth	Audit log (paginated, filterable)
GET	/api/v1/health	Shared	Health check
WS	/ws	Dashboard	WebSocket for real-time updates

AI Engine REST API (FastAPI)

Method	Endpoint	Description
POST	/api/v1/detect/anomaly	Isolation Forest anomaly detection
POST	/api/v1/explain	SHAP explanation for anomaly score
POST	/api/v1/classify/threat	Optional threat classification
GET	/api/v1/health	Health check

This HLD is governed by [ADR-001](file:///d:/AI%20SIEM/docs/architecture/decisions/ADR-001-modular-monolith.md) and implements the requirements defined in [SRS-001](file:///d:/AI%20SIEM/docs/SRS.md).

Document Version: 1.0
Last Updated: 2026-07-18